

## 2. Number Recognition with Keras

May 4, 2024

“You can think of solving MNIST as the”Hello world” of deep learning.” - François Chollet

The MNIST dataset consists of a collection of 70,000 images of handwritten numerical digits. Each image consists of a 28x28 pixel grayscale image.


The cell below displays 10 randomly selected images out of the training dataset.

```
[2]: import matplotlib.pyplot as plt
import numpy as np

# Load the MNIST dataset
from tensorflow.keras.datasets import mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

# Select a random sample of images
num_samples = 10
random_indices = np.random.choice(len(train_images), num_samples)

# Plot the random samples
for i, idx in enumerate(random_indices):
    plt.subplot(2, 5, i+1)
    plt.imshow(train_images[idx], cmap='gray')
    plt.title(f"Label: {train_labels[idx]}")
plt.tight_layout()
plt.show()
```



We can verify as follows that each image is simply a 2d numpy array, with grayscale intensity in the range 0 to 255.

```
[56]: example_image = train_images[random_indices[0]]
print(
    type(example_image),
    example_image,
    f"Number of Train images: {len(train_images)}",
    f"Number of Test images: {len(test_images)}",
    sep='\n'
)
```

```
<class 'numpy.ndarray'>
[[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  18  43 123 201 227 148 131  29  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  84 216 253 252 252 252 252 253 231  95  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  73 246 252 253 252 252 247 249 245 252 242 130
   0  0  0  0  0  0  0  0  0  0]
```

```

[ 0 0 0 0 0 0 127 252 252 190 189 136 63 70 49 211 252 231
 28 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 21 86 42 0 0 0 0 0 0 146 252 252
129 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 103 253 253
147 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 190 252 252
68 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 101 247 252 194
4 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 32 134 169 253 252 252 231
81 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 36 223 252 252 253 252 252 252
252 45 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 129 253 253 253 247 211 211 239
253 237 55 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 21 159 203 168 53 0 0 69
239 253 196 7 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 9 0 0 0 0 0
106 253 252 74 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 253 252 126 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 253 252 126 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 11 87 80 0 0 0 0 0 0 0 38
192 255 253 109 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 143 252 211 0 0 0 0 0 0 87 206
252 253 189 14 0 0 0 0 0 0]
[ 0 0 0 0 0 0 9 204 252 250 118 22 22 22 128 206 251 252
252 170 32 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 83 242 253 252 252 252 252 253 252 252 210
145 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 35 165 252 252 252 252 253 252 146 14
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0]]

```

Number of Train images: 60000

Number of Test images: 10000

# 1 Defining the Network Architecture

```
[57]: from tensorflow import keras
      from tensorflow.keras import layers

      model = keras.Sequential([
          layers.Dense(512, activation="relu"),
          layers.Dense(10, activation="softmax")
      ])
```

## 1.1 The Compilation Step

```
[58]: model.compile(optimizer="rmsprop",
                    loss="sparse_categorical_crossentropy",
                    metrics=["accuracy"])
```

## 1.2 Preparing the Training Data

The reshape we are doing can be represented as: 60000, 28, 28 - > 60000, 28\*28

```
[59]: train_images = train_images.reshape((60000, 28 * 28))
      train_images = train_images.astype("float32") / 255
      test_images = test_images.reshape((10000, 28 * 28))
      test_images = test_images.astype("float32") / 255
```

## 1.3 Fitting the Model

```
[60]: model.fit(train_images, train_labels, epochs=5, batch_size=128)
```

Epoch 1/5

469/469 [=====] - 6s 12ms/step - loss: 0.2645 - accuracy: 0.9235

Epoch 2/5

469/469 [=====] - 7s 14ms/step - loss: 0.1053 - accuracy: 0.9685

Epoch 3/5

469/469 [=====] - 7s 14ms/step - loss: 0.0701 - accuracy: 0.9791

Epoch 4/5

469/469 [=====] - 7s 14ms/step - loss: 0.0503 - accuracy: 0.9849

Epoch 5/5

469/469 [=====] - 7s 15ms/step - loss: 0.0376 - accuracy: 0.9890

```
[60]: <keras.callbacks.History at 0x19bde2710>
```

```
[61]: model.summary()
```

```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 512)	401920
dense_7 (Dense)	(None, 10)	5130

=====  
Total params: 407,050  
Trainable params: 407,050  
Non-trainable params: 0  
=====

## 1.4 Making Predictions on the Test Set

Each number of index `i` in that array corresponds to the probability that digit image `test_digits[0]` belongs to class `i`.

```
[62]: test_digits = test_images[0:10]
      predictions = model.predict(test_digits)
      predictions[0]
```

```
1/1 [=====] - 0s 103ms/step
```

```
[62]: array([1.9355136e-07, 7.8654594e-09, 6.5878298e-06, 1.2391029e-04,
            4.3995693e-12, 8.3856868e-08, 1.7252054e-11, 9.9986339e-01,
            4.2469955e-07, 5.3094614e-06], dtype=float32)
```

## 1.5 Evaluating the Model on the Test Set

```
[63]: test_loss, test_acc = model.evaluate(test_images, test_labels)
      print(f"test_acc: {test_acc}")
```

```
313/313 [=====] - 2s 6ms/step - loss: 0.0581 -
accuracy: 0.9808
test_acc: 0.9807999730110168
```

The test-set accuracy turns out to be 98.0%—that’s quite a bit lower than the training-set accuracy (98.9%). This gap between training accuracy and test accuracy is an example of overfitting: the fact that machine learning models tend to perform worse on new data than on their training data. Overfitting is a central topic in chapter 3.

Let’s plot some images from the test set along with the predictions by the model:

```
[65]: # Select a random sample of images
      num_samples = 10
```

```

random_indices = np.random.choice(len(test_images), num_samples)

# Plot the random samples with their predicted labels
for i, idx in enumerate(random_indices):
    test_prediction = model.predict(test_images[idx:idx+1])
    predicted_label = np.argmax(test_prediction, axis=1)
    plt.subplot(2, 5, i+1)
    plt.imshow(test_images[idx].reshape(28, 28), cmap='gray')
    plt.title(f"True: {test_labels[idx]}\nPred: {predicted_label[0]}")
plt.tight_layout()
plt.show()

```

```

1/1 [=====] - 0s 57ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 91ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 28ms/step

```

